COURSE: PHARMACEUTICAL BIOTECHNOLOGY & ADVANCED INFORMATICS

Artificial Intelligence in Drug Discovery and Repurposing

A Mini Research Project on Paradigm Shifts, Computational Methodologies, and Clinical Horizons

---

Topic Focus:	AI & Biotechnology in Pharma
Document Type:	Mini Research Project
Date:	May 24, 2026
Status:	Final Submission

Abstract

The conventional drug discovery pipeline is characterized by staggering financial costs, high attrition rates, and prolonged development timelines, often spanning over a decade. Recently, Artificial Intelligence (AI), driven by deep learning, generative chemistry, and high-throughput multi-omics data integration, has emerged as a disruptive paradigm capable of accelerating both *de novo* drug design and drug repurposing. This mini-research project investigates the core structural frameworks, algorithmic architectures, and clinical applications of AI in modern pharmacology. We review current deep learning methodologies, explore successful case studies of AI-repurposed compounds, examine mathematical evaluation metrics for binding affinity, and outline the lingering logistical, computational, and regulatory challenges that restrict seamless bench-to-bedside translation.

1. Introduction

Bringing a new molecular entity (NME) to market is historically inefficient. On average, the development of a single successful therapeutic agent costs between $1.3 billion and $2.8 billion, requiring 10 to 12 years from initial target validation to final regulatory authorization. Strikingly, over 90% of therapeutic candidates entering Phase I clinical trials fail due to unpredicted toxicity, insufficient clinical efficacy, or poor pharmacokinetic properties. This operational bottleneck has forced the pharmaceutical sector to explore computational alternatives capable of navigating the immense chemical space, which is estimated to contain up to 10^{60} synthesizable, drug-like molecules.

Artificial Intelligence (AI), specifically machine learning (ML) and deep learning (DL), provides an advanced toolset to address these structural deficits. By extracting non-linear feature relationships from colossal datasets—such as genomic sequences, crystal structures, phenotypic screens, and electronic health records (EHRs)—AI models can rapidly predict target binding, synthesize novel chemical scaffolds, and forecast toxicological outcomes. Furthermore, AI has revolutionized drug repurposing (repositioning), an economical strategy that identifies novel clinical indications for existing FDA-approved drugs, mitigating early-stage safety risks and reducing development timelines to 3–5 years.

This paper presents an analytical exploration of AI's functional roles within modern pharmacology. Section 2 establishes the baseline workflow of traditional versus AI-accelerated pipelines. Section 3 delineates primary machine learning techniques, covering deep generative models and graph neural networks. Section 4 evaluates the technical mechanisms of AI-driven drug repurposing. Section 5 discusses mathematical modeling frameworks, and Section 6 presents case studies, followed by an objective examination of current implementation barriers and future perspectives.

2. The Pharmaceutical Paradigm Shift

To understand the utility of AI, it is essential to contextualize the stages of traditional drug discovery against an AI-augmented workflow. The classical framework proceeds linearly through target identification, lead generation, lead optimization, preclinical in vitro/in vivo testing, and clinical trial evaluation. AI compresses the initial stages through high-throughput virtual screening, predictive ADME-Tox (Absorption, Distribution, Metabolism, Excretion, and Toxicity) modeling, and generative chemical architecture design.

Table 1: Comparative breakdown of classical pharmaceutical workflows versus AI-integrated implementation frameworks.

Pipeline Phase	Traditional Methodologies	AI-Augmented Paradigm	Time Impact Reduction
Target Identification	Literature review, manual cellular assays, phenotypic screening.	Multi-omics data mining, network pharmacology, semantic knowledge graphs.	65% time reduction
Lead Generation & Screening	Physical High-Throughput Screening (HTS) of physical chemical libraries.	Deep-learning virtual screening, ensemble docking, active learning.	80% time reduction
Lead Optimization	Iterative medicinal chemistry synthesis, SAR experimentation.	Generative Adversarial Networks (GANs), Reinforcement Learning (RL).	70% time reduction
Preclinical ADME-Tox	Animal model parsing, extensive in vitro toxicity screens.	Deep neural net quantitative structure-activity relationship (QSAR) models.	50% cost reduction

3. Core AI Technologies in De Novo Drug Design

De novo drug design refers to the generation of completely novel chemical entities optimized to modulate specific biological targets without utilizing known active starting scaffolds. AI architectures excel in this space by treating molecular design either as a sequence-generation problem or as a non-Euclidean graph topology optimization challenge.

3.1 Generative Chemistry and Language Models

Molecules are frequently represented computationally using the Simplified Molecular Input Line Entry System (SMILES), a notation system that maps molecular graphs into linear textual strings. For instance, benzene is written as `C1=CC=CC=C1`. By leveraging this linear representation, deep learning architectures originally engineered for Natural Language Processing (NLP)—such as Recurrent Neural Networks (RNNs), Long Short-

Term Memory (LSTM) networks, and Transformers—can be trained on massive molecular datasets (e.g., ChEMBL, PubChem) to learn structural grammar.

Once trained, Generative Adversarial Networks (GANs) and Variational Autoencoders (VAEs) manipulate latent space configurations to sample new chemical structures. In a VAE framework, an encoder compresses discrete SMILES text into a continuous latent vector space, and a decoder attempts to reconstruct valid chemical strings from that space. By applying Reinforcement Learning (RL) policy gradients, the generator is penalized or rewarded based on predefined molecular objectives, such as binding affinity coefficients, synthetic accessibility (SA) scores, or water-octanol partition coefficients ($\log P$).

3.2 Graph Neural Networks (GNNs)

While linear strings are highly computable, they fail to capture the multi-dimensional spatial arrangements and spatial geometric nuances of molecules. Graph Neural Networks (GNNs), including Graph Convolutional Networks (GCNs) and Message Passing Neural Networks (MPNNs), overcome this limitation by representing molecules directly as mathematical graphs $G = (V, E)$, where V represents the vertices (atoms characterized by element type, hybridization, and formal charge) and E denotes the edges (chemical bonds characterized by bond order and stereochemistry). GNNs pass vector messages across neighboring atoms to generate continuous structural embeddings, vastly improving the accuracy of quantitative structure-property relationship (QSPR) predictions and structural activity evaluations.

4. Methodologies in AI-Driven Drug Repurposing

Drug repurposing leverages existing safety profiles of approved compounds to identify secondary therapeutic indications, circumventing preclinical toxicity testing. AI optimizes this pathway using two main approaches: network-based pharmacology and transcriptomic deep embedding.

4.1 Network Pharmacology and Knowledge Graphs

Biomedical knowledge graphs (KGs) organize heterogeneous biomedical information by modeling biological entities (e.g., diseases, genes, proteins, pathways, and approved compounds) as nodes, and their empirical associations as directed edges. AI utilizes graph embedding techniques, such as Node2Vec or Graph Attention Networks (GATs), to project these complex relational webs into low-dimensional vector spaces. Link prediction algorithms can then mathematically calculate the probability that an edge should exist between a specific drug node and a disease node, identifying potential therapeutic applications that are not apparent through traditional isolated biochemical screening methods.

4.2 Transcriptomic Alignment and Phenotypic Signatures

Another dominant paradigm relies on matching transcriptomic perturbation signatures. Deep autoencoders are trained on cellular expression datasets, such as the Connectivity Map (CMap) or the Library of Integrated

Network-Based Cellular Signatures (LINCS). When a disease displays a distinct up-regulated or down-regulated gene expression profile, AI algorithms screen a database of drug transcriptomic signatures to identify molecules that produce an inverse profile. If a disease up-regulates a pathogenic pathway, the model identifies a drug that down-regulates that exact pathway, indicating a functional therapeutic offset.

5. Mathematical Frameworks & Evaluation Metrics

Evaluating AI models in pharmacology requires rigorous quantitative metrics. In virtual screening, predictive tasks are often framed as regression problems estimating the binding affinity (K_i or IC_{50}) or binary classification problems predicting active versus inactive status. A core objective function for predicting free energy of binding (ΔG_{bind}) using deep neural networks can be formulated as an optimization over empirical loss functions.

Let y_i represent the experimentally derived binding affinity, and $f(x_i; \theta)$ denote the predictive value outputted by an AI network parameterised by weights θ given structural input features x_i. The objective function often minimizes the Mean Squared Error (MSE) paired with an L_2 regularization constraint to prevent overfitting:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^{N} (y_i - f(x_i; \theta))^2 + \lambda \|\theta\|_2^2$$

where λ controls the penalty on weight magnitude to maintain structural generalization capabilities across unseen chemical domains.

To quantify the accuracy of generative distribution models, the Fréchet ChemNet Distance (FCD) is routinely calculated. The FCD measures the distance between the distribution of generated molecules (p_g) and real active compounds (p_r) by passing them through the hidden layers of a specialized neural network trained on chemical properties (ChemNet):

$$FCD = \|\mu_r - \mu_g\|_2^2 + Tr(\Sigma_r + \Sigma_g - 2(\Sigma_r \Sigma_g)^{0.5})$$

where μ represents the mean vector of the layer activations, Σ denotes the covariance matrix, and Tr signifies the trace of the resulting matrix operator. A lower FCD indicates that the generated molecules align closely with the chemical space of verified therapeutics.

6. Empirical Case Studies

The practical utility of AI in drug discovery is demonstrated by several landmark compounds that have advanced into clinical evaluation.

6.1 DSP-1181 (Obsessive-Compulsive Disorder)

Created via a joint venture between Sumitomo Dainippon Pharma and Exscientia, DSP-1181 became the first fully AI-designed molecule to enter human clinical trials (Phase I) in 2020. The compound acts as a long-acting serotonin 5-HT1A receptor agonist. While traditional exploratory chemistry stages require an average of 4.5 years, the AI platform completed the exploratory discovery phase in under 12 months by generating and evaluating millions of structural variations to optimize target selectivity and metabolic stability simultaneously.

6.2 INS018_055 (Idiopathic Pulmonary Fibrosis)

Developed by Insilico Medicine, INS018_055 is a small-molecule inhibitor targeting idiopathic pulmonary fibrosis (IPF), a chronic, progressive lung disease. Using generative chemistry models (Generative Adversarial Networks) combined with reinforcement learning, Insilico identified a novel biological target and synthesized a highly selective inhibitor. In 2023, this candidate advanced into Phase II clinical trials, representing a critical milestone where both target selection and molecular design were completely directed by AI engines.

6.3 Baricitinib Repurposing for COVID-19

During the early stages of the SARS-CoV-2 pandemic, BenevolentAI utilized its automated knowledge graph platform to identify approved therapeutics that could inhibit the viral entry mechanism of SARS-CoV-2 while dampening the life-threatening systemic hyper-inflammatory cascade (cytokine storm). Within 48 hours, the platform identified Baricitinib—an existing Janus Kinase (JAK) inhibitor approved for rheumatoid arthritis. Subsequent clinical validations confirmed its efficacy, culminating in an Emergency Use Authorization (EUA) and final FDA approval for severe COVID-19 cases, demonstrating the rapid turnaround times of network-driven repurposing models.

7. Technical Challenges and Translational Barriers

Despite these technical milestones, integrating AI into the pharmaceutical industry presents several operational challenges. Addressing these bottlenecks is necessary to ensure AI-driven candidates routinely succeed in later clinical development stages.

- **Data Sparsity, Imbalance, and Quality:** Deep learning models require massive, structured datasets. However, proprietary academic and corporate silos isolate high-quality assay datasets. Furthermore, published scientific literature contains profound publication bias, consistently presenting positive results while omitting negative screening outcomes. Training AI models on positive data alone leads to high false-positive generation rates.
- **The "Black Box" Problem and Interpretability:** Deep neural networks often lack mechanistic transparency. Medicinal chemists and regulatory authorities (such as the FDA and EMA) require clear causal explanations for why an AI algorithm predicts a specific molecule will be non-toxic or highly active. Without clear explainability frameworks (e.g., SHAP, Integrated Gradients), establishing safety profiles remains challenging.
- **Chemical Synthesis Constraints:** Early generative architectures designed molecular graphs that achieved high binding scores in virtual environments but were structurally impossible to synthesize in physical laboratories due to steric hindrance or unstable chemical rings. While current models integrate synthetic accessibility filters, aligning generative outputs with automated robotic synthesis workflows remains an active area of optimization.

8. Future Perspectives and Conclusion

The integration of AI with quantum computing and automated synthesis platforms represents the next frontier in digital pharmacology. Quantum machine learning algorithms possess the theoretical capacity to solve the Schrödinger equation for multi-electron molecular configurations with absolute precision. This advancement will enable accurate prediction of molecular electronic properties and binding affinities, minimizing reliance on empirical approximations.

Simultaneously, the development of closed-loop "Design-Make-Test-Analyze" (DMTA) autonomous laboratories is increasing. In these facilities, an AI generator designs a chemical scaffold, transmits the synthetic instructions to robotic fluidic synthesizers, executes automated microfluidic bioassays, and feeds the resulting functional data back into its central learning loop to refine the next generation of compounds without human intervention.

In conclusion, AI is reshaping the pharmaceutical industry, shifting it from an empirical, slow discovery paradigm to an agile, data-driven framework. While data silos and interpretability challenges remain, successful clinical entries like INS018_055 and repurposed discoveries like Baricitinib demonstrate that AI can reduce timelines, lower operational costs, and identify novel therapeutic options. As computational models improve and integrate with automated laboratories, AI will become a core element of modern drug discovery, accelerating the delivery of life-saving therapeutics to patients worldwide.

References

1. Bender, A., & Cortés-Ciriano, I. (2021). Artificial intelligence in drug discovery: what is realistic, what are the illusions? *Drug Discovery Today*, 26(2), 511-524.

2. DiMasi, J. A., Grabowski, H. G., & Hansen, R. W. (2016). Innovation in the pharmaceutical industry: New estimates of R&D costs. *Journal of Health Economics*, 47, 20-33.

3. Pushpakom, S., Iorio, F., Eyers, P. A., et al. (2019). Drug repurposing: progress, challenges and recommendations. *Nature Reviews Drug Discovery*, 18(1), 41-58.

4. Senior, A. W., Evans, R., Jumper, J., et al. (2020). Improved protein structure prediction using potentials from deep learning. *Nature*, 577(7792), 706-710.

5. Stokes, J. M., Yang, K., Swanson, K., et al. (2020). A deep learning approach to antibiotic discovery. *Cell*, 180(4), 688-702.

6. Stebbing, J., Sánchez-Nievas, G., Falcone, M., et al. (2020). JAK inhibition with baricitinib suppresses treatment-resistant COVID-19. *EMBO Molecular Medicine*, 12(11), e13007.

7. Zhavoronkov, A., Ivanenkov, Y. A., Aliper, A., et al. (2019). Deep learning enables rapid identification of potent DDR1 kinase inhibitors. *Nature Biotechnology*, 37(9), 1038-1040.